

Reduced Matrix Multiplication Report

1. Srinjoy Mukherjee(SrNo: 21152)
2. Adesh Adikane(SrNo: 21027)

The link for the github repository: [Reduced Matrix Multiplication](#)

Part A

In this part of the assignment, we are supposed to implement single threaded and multi-threaded Reduced Matrix Multiplication(RMM) . Naive implementation is given for reference.Improving the performance of the operation can be achieved using many ways. In the previous assignment , we explored improvement by using blocked matrix multiplication and loop interchange.We observed significant improvement using those techniques. In this assignment we use vector instructions to exploit data parallelism. We make a avx vector of 256 bits/4 integers and do the addition and multiplication operations with the corresponding vectors of the other matrix.We also use loop interchange to utilise locality of reference.We observe significant improvement over naive implementation of RMM.The reason for the improvement is using data parallelism, SIMD instructions, where a single instruction is fetched and decoded but executed on multiple elements.In multi-threaded implementation we use 8 threads . Each thread is responsible for calculating RMM for $(N/8)$ number of rows of resultant matrix where N is the size of resultant matrix.

Bottleneck: In single threaded we can't exploit much higher parallelism as one thread may be blocked or waiting for cache, so we move on to multi-threading in which multiple threads can work on different parts of the matrix, which will result in higher performance.


The code was run on CLSERV server .The specs of the machine are as follows:

```
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Byte Order:        Little Endian
CPU(s):            32
On-line CPU(s) list: 0-31
Thread(s) per core: 2
Core(s) per socket: 8
Socket(s):         2
NUMA node(s):      2
Vendor ID:         GenuineIntel
CPU family:        6
Model:             63
Model name:        Intel(R) Xeon(R) CPU E5-2640 v3 @ 2.60GHz
Stepping:          2
CPU MHz:           1199.219
CPU max MHz:       3400.0000
CPU min MHz:       1200.0000
BogoMIPS:          5194.02
Virtualization:    VT-x
L1d cache:         32K
L1i cache:         32K
L2 cache:          256K
L3 cache:          20480K
```


Charts

The charts for different implementations for each matrix size:


1. Time(in ms):




2. L1 Read Cahce Misses(in logarithm scale):




3. LL Read Cache Misses(in logarithm scale):




4. LL Write Cache Misses(in logarithm scale):



5. TLB Misses(in logarithm scale):



6. Page Faults(in logarithm scale):



Observations

We make the graphs of the different counters we used to compare the performance. We compare the performance on the following parameters:

- 1. Time
- 2. L1 Read Miss
- 3. LL Read Miss
- 4. LL Write Miss
- 5. TLB Miss
- 6. Page Faults

Observation Table

IMPLEMENTATION	SIZE	TIME	L1 READ MISS	LL READ MISS	LL WRITE MISS	TLB MISS	PAGE FAULT
REFERENCE	16	0.107	37	0	0	1	0
SINGLE-THREADED	16	0.08	43	0	0	0	0
MULTI-THREADED	16	0.831	1460	64	10	179	16
REFERENCE	4096	986309	30981154829	18278461380	311788	17350628081	0
SINGLE-THREADED	4096	258387	11289280138	5115642215	179371	6813772779	40917
MULTI-THREADED	4096	42577.9	1830	392	91	373	30
REFERENCE	8192	8539590	2.96909E+11	1.33618E+11	1529653	1.39737E+11	0
SINGLE-THREADED	8192	2549420	1.1181E+11	52095504886	3581961	55496798295	32768
MULTI-THREADED	8192	421384	1667	429	79	420	38
REFERENCE	16384	69689700	3.06613E+12	1.85347E+12	8174646	1.12464E+12	8337199

IMPLEMENTATION	SIZE	TIME	L1 READ MISS	LL READ MISS	LL WRITE MISS	TLB MISS	PAGE FAULT
SINGLE-THREADED	16384	23426700	9.64299E+11	5.60175E+11	16979372	4.64353E+11	419786
MULTI-THREADED	16384	3457470	1805	594	108	495	42

Looking at the table, we observe that the numbers are too high, so for ease of plotting them on the graph, we simply take the performance improvement ratio. For each parameter it is calculated by: (Reference Value/Improved Value).

Improvement Ratio Table

IMPLEMENTATION	SIZE	TIME
REFERENCE	16	1
SINGLE-THREADED	16	1.3375
MULTI-THREADED	16	0.128
REFERENCE	4096	1
SINGLE-THREADED	4096	3.817
MULTI-THREADED	4096	23.16
REFERENCE	8192	1
SINGLE-THREADED	8192	3.349
MULTI-THREADED	8192	20.265
REFERENCE	16384	1
SINGLE-THREADED	16384	2.9747
MULTI-THREADED	16384	20.15

As we look into the table, we observe that for very small matrix sizes, the performance for naive implementation is better than single and multi-threaded. This is because creating the vector variables from data in matrix takes significant percentage of entire execution time. As we keep increasing the size, the usage of vector instructions becomes more prominent. Among the three implementations, the multi-threaded performed the best as expected. It performs the best because firstly, it uses loop interchange utilising locality of reference, secondly vector instructions to exploit data parallelism and finally multi-threading which divides the tasks among multiple threads and which can be processed parallelly.

Part B

Implementation:-

1. We are allocating the memory for each input matrix and the output matrix on the GPU device.
2. Multiplication function computes the matrix multiplication on the GPU
3. We are computing by considering different cases i.e by taking different ThreadBlock sizes and and no of ThreadBlocks.
4. In GPU the hierarchy is that there is a grid which contains no of ThreadBlocks and each ThreadBlock contains multiple no of threads. ThreadBlocks run in Parallel on each SM(Simultaneous Multiprocessor).


- 5. Here we are taking one grid having 2 dimensions x and y each dimension contains no of thread blocks. We are taking dimension of grid by dividing the output matrix size with threadblock size. Each thread will compute each element of output matrix.
- 6. Each time we are taking ThreadBlock of different sizes (16,32,1024,etc).


The code was run on NVidia 3050. the specs are as below:

```
Device 0: "NVIDIA GeForce RTX 3050 Laptop GPU"
CUDA Driver Version / Runtime Version      12.0 / 11.8
CUDA Capability Major/Minor version number: 8.6
Total amount of global memory:              4096 MBytes (4294443008 bytes)
(016) Multiprocessors, (128) CUDA Cores/MP: 2048 CUDA Cores
GPU Max Clock rate:                        1500 MHz (1.50 GHz)
Memory Clock rate:                         5871 Mhz
Memory Bus Width:                          128-bit
Warp size:                                 32
Maximum number of threads per multiprocessor: 1536
Maximum number of threads per block:        1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
```


Charts

Charts of time(in ms) for different block sizes for matrix size:






2. 4096:



3. 8192:



4. 16384:

Observations:-

Table for Execution Time(in ms) of Block Size vs Matrix Size.

Block Size	512*512	1024*1024	2048*2048	4096*4096	8192*8192	16384
4	0.899235	9.929249	74.031097	490.072510	4094.792969	34107.718750
8	0.46016	2.886912	25.359489	191.907166	1957.659180	14709.500000
16	0.492256	1.398688	24.898174	180.233398	1352..780273	11524.953125
32	0.407872	3.123168	25.817345	189.204437	1488.171875	12057.994141

Table for maximum Speedup achieved for each matrix size.

Matrix Size	Reference Time	GPU Time	SpeedUp
512 x 512	750 ms	0.492256 ms	1523.77
1024 x 1024	6496 ms	1.398688 ms	4644.6
2048 x 2048	104961 ms	24.898174 ms	4217.5
4096 x 4096	986309 ms	180.233398 ms	5472.4
8192 x 8192	55395962 ms	1352.780273 ms	4092.28
16384 x 16384	69683723 ms	11524.953125 ms	6046.33

1. Increasing the Threadblock sizes reduces the execution time of the reduced matrix multiplication.
2. Threadblock size of 16 gives minimum execution time as compared to other block sizes.
3. In case of ThreadBlock size 16 we are utilizing the resources (SM and Threads in each SM) with full capacity as compared to other other block size.
4. We can see that with small blocksize we are not using the SM (Simultaneous Multiprocessor. with full capacity.
5. We can also see that with large blocksize we are not using enough SM i.e we are not using enough parallelism.
6. To get the best possible execution time there needs to be good balance between the no of Threadsblocks and size of the thread block.
7. We cannot see much difference in the execution time with less matrix size. For less matrix size most of the time is taken by memory allocation and transferring of data on GPU device, as we increase the matrix size we can clearly see the difference in execution time and speedUp.

BottleNeck:- Since there are limited number of cores in CPU,, so we are moving to GPU.As GPU can have maximum ThreadBlock size of 1024 threads so for matrix size greater than that we need to increase no. of blocks.

Conclusion

For PartA, We observed the performance improvement due to AVX Instructions and reasoned about it .We also saw the usefulness of multi-threaded program and how it affects the execution time.So we can conclude that exploiting data parallelism and instruction parallelism gives us significant performance gain. For PartB , as GPU contains a huge no of cores which can process a lot of data , thus the throughput is considerably higher, although for lower data sizes, it performs worse as the bottleneck of copying data from CPU to GPU becomes significant percentage of the entire program., hence the worse performance.